

## bcc: Taming Linux 4.3+ Tracing Superpowers

22 Sep 2015

Here's a quick tour of some new open source tools which I demonstrated at the Silicon Valley Linux Technology meetup last night. These use the new eBPF capabilities added in recent Linux, including Linux 4.3.

Summarizing the distribution of disk I/O latency:

```
# ./biolateness
Tracing block device I/O... Hit Ctrl-C to end.
^C
```

usecs	: count	distribution
0 -> 1	: 0	
2 -> 3	: 0	
4 -> 7	: 0	
8 -> 15	: 0	
16 -> 31	: 0	
32 -> 63	: 0	
64 -> 127	: 1	
128 -> 255	: 12	*****
256 -> 511	: 15	*****
512 -> 1023	: 43	*****
1024 -> 2047	: 52	*****
2048 -> 4095	: 47	*****
4096 -> 8191	: 52	*****
8192 -> 16383	: 36	*****
16384 -> 32767	: 15	*****
32768 -> 65535	: 2	*
65536 -> 131071	: 2	*

Tracing per-disk I/O:

```
# ./biosnoop
TIME(s)      COMM          PID    DISK    T  SECTOR  BYTES  LAT(ms)
0.000004001  supervise      1950    xvda1   W  13092560 4096   0.74
0.000178002  supervise      1950    xvda1   W  13092432 4096   0.61
0.001469001  supervise      1956    xvda1   W  13092440 4096   1.24
0.001588002  supervise      1956    xvda1   W  13115128 4096   1.09
1.022346001  supervise      1950    xvda1   W  13115272 4096   0.98
1.022568002  supervise      1950    xvda1   W  13188496 4096   0.93
1.023534000  supervise      1956    xvda1   W  13188520 4096   0.79
1.023585003  supervise      1956    xvda1   W  13189512 4096   0.60
```

Tracing the open() syscall:

```
# ./opensnoop
PID  COMM          FD  ERR  PATH
17326 <...>          7    0  /sys/kernel/debug/tracing/trace_pipe
17358 run            3    0  /lib/x86_64-linux-gnu/libtinfo.so.5
17358 run            3    0  /lib/x86_64-linux-gnu/libdl.so.2
17358 run            3    0  /lib/x86_64-linux-gnu/libc.so.6
17358 run          -1    6  /dev/tty
17358 run            3    0  /proc/meminfo
17358 run            3    0  /etc/nsswitch.conf
```

Counting VFS operation types:

```
# ./vfsstat
TIME      READ/s  WRITE/s  CREATE/s  OPEN/s  FSYNC/s
18:35:35:    241    15      4        99      0
18:35:36:    232    10      4        98      0
18:35:37:    244    10      4       107      0
18:35:38:    235    13      4        97      0
18:35:39:   6749  2633    4      1446      0
18:35:40:    277    31      4       115      0
```

Counting kernel function calls per-second, that match "tcp\*send\*":

```
# ./funccount -i 1 'tcp*send*'
Tracing... Ctrl-C to end.
```

ADDR	FUNC	COUNT
ffffffff816d2281	tcp_send_delayed_ack	30
ffffffff816d6c81	tcp_v4_send_check	31
ffffffff816c2f61	tcp_sendmsg	31
ffffffff816bf851	tcp_send_mss	31

ADDR	FUNC	COUNT
ffffffff816d1db1	tcp_send_fin	3
ffffffff816d0f71	tcp_send_ack	18
ffffffff816d2281	tcp_send_delayed_ack	214
ffffffff816c2f61	tcp_sendmsg	231
ffffffff816bf851	tcp_send_mss	231
ffffffff816d6c81	tcp_v4_send_check	255

ADDR	FUNC	COUNT
ffffffff816d0f71	tcp_send_ack	2
ffffffff816d2281	tcp_send_delayed_ack	9
ffffffff816c2f61	tcp_sendmsg	30
ffffffff816bf851	tcp_send_mss	30

Timing tcp\_sendmsg() latency (call duration), in microseconds:

```
# ./funclatency -u tcp_sendmsg
Tracing tcp_sendmsg... Hit Ctrl-C to end.
^C
```

usecs	count	distribution
0 -> 1	20778	*****
2 -> 3	15429	*****
4 -> 7	355	
8 -> 15	171	
16 -> 31	106	
32 -> 63	9	

Detaching...

... all of these tools have man pages, and most have help messages as well:

```
# ./funclatency -h
usage: funclatency [-h] [-p PID] [-i INTERVAL] [-T] [-u] [-m] [-r] pattern

Time kernel functions and print latency as a histogram

positional arguments:
  pattern                search expression for kernel functions

optional arguments:
  -h, --help            show this help message and exit
  -p PID, --pid PID     trace this PID only
  -i INTERVAL, --interval INTERVAL
                        summary interval, seconds
  -T, --timestamp       include timestamp on output
  -u, --microseconds    microsecond histogram
  -m, --milliseconds    millisecond histogram
  -r, --regex           use regular expressions. Default is "*" wildcards
                        only.

examples:
  ./funclatency do_sys_open      # time the do_sys_open() kernel function
  ./funclatency -u vfs_read     # time vfs_read(), in microseconds
  ./funclatency -m do_nanosleep # time do_nanosleep(), in milliseconds
  ./funclatency -mTi 5 vfs_read # output every 5 seconds, with timestamps
  ./funclatency -p 181 vfs_read # time process 181 only
  ./funclatency 'vfs_fstat*'    # time both vfs_fstat() and vfs_fstatat()
```

## Linux 4.3+, eBPF

What's new in Linux 4.3 is the ability to print strings from Extended Berkeley Packet Filters (eBPF) programs. This is only a small addition, but one I needed for many tools. eBPF is a virtual machine for running user-defined, sandboxed bytecode, with maps for data storage. I wrote about it in [eBPF: One Small Step](#).

eBPF enhances Linux tracing, allowing mini programs to be executed on tracing events. In my tools above, eBPF lets me tag events with custom timestamps, store histograms, filter events, and only emit summarized info to user-level. These capabilities give me the info I want, with the lowest possible overhead cost.

While eBPF provides amazing superpowers, there is a catch: it's hard to use via its assembly or C interface. The challenge attracts me, but it can be a brutal experience, especially if you write eBPF assembly directly (eg, see

bpf\_insn\_prog[] from [sock example.c](#); I've yet to code one of these from scratch that compiles). The C interface is better (see other examples in [samples/bpf](#)), but it's still laborious and difficult to use.

## Enter bcc

The BPF Compiler Collection ([bcc](#)) project provides a front-end for eBPF, making it easier to write programs. It uses C for the back-end instrumentation, and Python for the front-end interface. My tools at the start of this post all use bcc, and can be found in the [tools directory](#) of bcc on github. Also browse that directory for the `_example.txt` files.

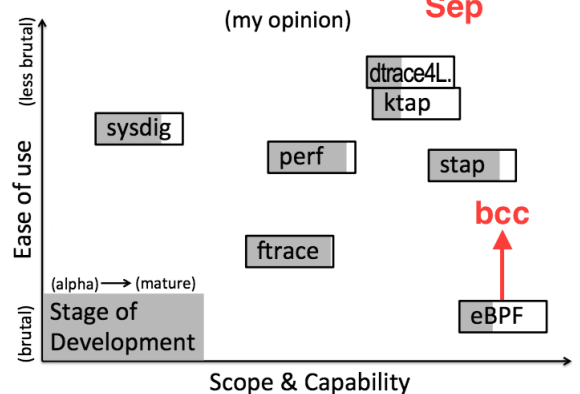
I've modified my diagram on the right (from [Velocity 2015](#)) to show the role that bcc plays: it improves the useability of eBPF.

It's still early days for bcc, and right now it isn't easy to setup and use, even once you're on Linux 4.3+ (eg, here are my own [notes](#)). In the future, this should be as simple as a package add.

And this will be adding user-level software only: the kernel parts (eBPF) are already in the Linux kernel.

Just as one example of bcc, here's the full code to my [biolatency tool](#), which I showed at the top of this post:

### The Tracing Landscape, ~~May~~ Sep 2015



```
1  #!/usr/bin/python
2  #
3  # biolatency    Summarize block device I/O latency as a histogram.
4  #                For Linux, uses BCC, eBPF.
5  #
6  # USAGE: biolatency [-h] [-T] [-Q] [-m] [interval] [count]
7  #
8  # Copyright (c) 2015 Brendan Gregg.
9  # Licensed under the Apache License, Version 2.0 (the "License")
10 #
11 # 20-Sep-2015   Brendan Gregg   Created this.
12
13 from __future__ import print_function
14 from bcc import BPF
15 from time import sleep, strftime
16 import argparse
17
18 # arguments
19 examples = """examples:
20     ./biolatency           # summarize block I/O latency as a histogram
21     ./biolatency 1 10      # print 1 second summaries, 10 times
22     ./biolatency -mT 1     # 1s summaries, milliseconds, and timestamps
23     ./biolatency -Q        # include OS queued time in I/O time
24 """
25 parser = argparse.ArgumentParser(
26     description="Summarize block device I/O latency as a histogram",
27     formatter_class=argparse.RawDescriptionHelpFormatter,
28     epilog=examples)
29 parser.add_argument("-T", "--timestamp", action="store_true",
30     help="include timestamp on output")
31 parser.add_argument("-Q", "--queued", action="store_true",
32     help="include OS queued time in I/O time")
```

```

33 parser.add_argument("-m", "--milliseconds", action="store_true",
34                     help="millisecond histogram")
35 parser.add_argument("interval", nargs="?", default=99999999,
36                     help="output interval, in seconds")
37 parser.add_argument("count", nargs="?", default=99999999,
38                     help="number of outputs")
39 args = parser.parse_args()
40 countdown = int(args.count)
41 debug = 0
42
43 # load BPF program
44 bpf_text = """
45 #include <uapi/linux/ptrace.h>
46 #include <linux/blkdev.h>
47
48 BPF_TABLE("array", int, u64, dist, 64);
49 BPF_HASH(start, struct request *);
50
51 // time block I/O
52 int trace_req_start(struct pt_regs *ctx, struct request *req)
53 {
54     u64 ts = bpf_ktime_get_ns();
55     start.update(&req, &ts);
56     return 0;
57 }
58
59 // output
60 int trace_req_completion(struct pt_regs *ctx, struct request *req)
61 {
62     u64 *tsp, delta;
63
64     // fetch timestamp and calculate delta
65     tsp = start.lookup(&req);
66     if (tsp == 0) {
67         return 0;        // missed issue
68     }
69     delta = bpf_ktime_get_ns() - *tsp;
70     FACTOR
71
72     // store as histogram
73     int index = bpf_log2l(delta);
74     u64 *leaf = dist.lookup(&index);
75     if (leaf) (*leaf)++;
76
77     start.delete(&req);
78     return 0;
79 }
80 """
81 if args.milliseconds:
82     bpf_text = bpf_text.replace('FACTOR', 'delta /= 1000000;')
83     label = "msecs"

```

```

84     else:
85         bpf_text = bpf_text.replace('FACTOR', 'delta /= 1000;')
86         label = "usecs"
87     if debug:
88         print(bpf_text)
89
90     # load BPF program
91     b = BPF(text=bpf_text)
92     if args.queued:
93         b.attach_kprobe(event="blk_account_io_start", fn_name="trace_req_start")
94     else:
95         b.attach_kprobe(event="blk_start_request", fn_name="trace_req_start")
96         b.attach_kprobe(event="blk_mq_start_request", fn_name="trace_req_start")
97     b.attach_kprobe(event="blk_account_io_completion",
98                     fn_name="trace_req_completion")
99
100    print("Tracing block device I/O... Hit Ctrl-C to end.")
101
102    # output
103    exiting = 0 if args.interval else 1
104    dist = b.get_table("dist")
105    while (1):
106        try:
107            sleep(int(args.interval))
108        except KeyboardInterrupt:
109            exiting=1
110
111        print()
112        if args.timestamp:
113            print("%-8s\n" % strftime("%H:%M:%S"), end="")
114
115        dist.print_log2_hist(label)
116        dist.clear()
117
118        countdown -= 1
119        if exiting or countdown == 0:
120            exit()

```

biolatenacy hosted with ❤ by GitHub

[view raw](#)

A lot of this code is logic for processing command-line arguments. The C instrumentation code is defined in-line, as `bpf_text`. There are a few things I'd like improved, like switching to tracepoints instead of kprobes, but this is not bad so far. I can write tools in this.

`bcc` can also do a lot more than my examples here: it can also be used for advanced network traffic control. For more about those capabilities, see the [IO Visor](#) project.

Even if few people learn `bcc` programming, it should see success through the use of its tools. At a company like Netflix, it will only take a few of us learning `bcc`/eBPF to have a big impact for the company: we can develop tools for others to use, and plugins for our other analysis software (eg, [Vector](#)). A lot of what we create is open sourced, so others will benefit as well.

# Other tracers and tools

bcc won't be the only interface to eBPF. There is already work on bringing it to Linux perf\_events. It would also be great to see a tracer with a grammar, like SystemTap or ktap, support eBPF, which would make ad hoc tools even easier to write. (Perhaps bcc will provide its own grammar in the future.) eBPF should make other enhancements possible for the other tracers.

I previously created [perf-tools](#), a collection of mostly ftrace-based tracing tools for Linux systems. With eBPF and bcc, I'll eventually switch some of these tools to eBPF, where they will have more features, be easier to maintain, and have lower overhead. For example, my perf-tools [iolatency](#) tool, which is equivalent to biolatency, processed every disk event in user-level, costing measurable overhead (which I warned about in the tool and its man page). The overhead for the bcc biolatency version should be negligible.

With bcc/eBPF, I'll also be creating many new tools that were previously impossible (or impractical) to do.

*Thanks to Alexei Starovoitov and Brenden Blanco of PLUMgrid, and others, for developing eBPF and bcc, and Deirdré Straughan for edits to this post.*

---

*You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).*