

## execsnoop For Linux: See Short-Lived Processes

28 Jul 2014

Every time I can't immediately explain CPU usage, I wonder if short-lived processes are to blame. On Linux systems I can debug this using [atop](#), which uses process accounting to catch these fleeting processes. But I wish I had my execsnoop tool, which creates a live log of each process for later study.

I just ported it to Linux. Here's an example, where "man ls" is run in another window:

```
# ./execsnoop
Tracing exec()s. Ctrl-C to end.
  PID   PPID  ARGS
20139   20135 mawk -W interactive -v o=1 -v opt_name=0 -v name= [...]
20140   20138 cat -v trace_pipe
20171   16743 man ls
20178   20171 preconv -e UTF-8
20181   20171 pager -s
20180   20171 nroff -mandoc -rLL=173n -rLT=173n -Tutf8
20179   20171 tbl
20184   20183 locale charmap
20185   20180 groff -mtty-char -Tutf8 -mandoc -rLL=173n -rLT=173n
20186   20185 troff -mtty-char -mandoc -rLL=173n -rLT=173n -Tutf8
20187   20185 grotty
```

Great! The first two lines, showing mawk and cat, are from execsnoop initializing. The remaining show the arcane workings of the man command.

Here's catching part of a Linux build:

```
# ./execsnoop
Tracing exec(s). Ctrl-C to end.
  PID  PPID  ARGS
25753  25588 /bin/sh -c echo 'python' | grep ^/ -q && echo y
25755  25753 grep ^/ -q
25756  25588 sh -c command -v 'python' | awk 'NR==1 {t=$0} NR>1 {t=t "m822df3020w6a44id34bt574ct
25758  25756 awk NR==1 {t=$0} NR>1 {t=t "m822df3020w6a44id34bt574ctac44eb9f4n" $0} END {printf t
25759  25588 /bin/sh -c echo 'python' | grep ^/ -q && echo y
25761  25759 grep ^/ -q
25762  25588 sh -c command -v 'python' | awk 'NR==1 {t=$0} NR>1 {t=t "m822df3020w6a44id34bt574ct
25764  25762 awk NR==1 {t=$0} NR>1 {t=t "m822df3020w6a44id34bt574ctac44eb9f4n" $0} END {printf t
25765  25588 /bin/sh -c echo '/usr/bin/python-config' | grep ^/ -q && echo y
25767  25765 grep ^/ -q
25768  25588 sh -c test -f '/usr/bin/python-config' -a -x '/usr/bin/python-config' && echo y
25769  25588 /bin/sh -c echo '/usr/bin/python-config' | grep ^/ -q && echo y
25771  25769 grep ^/ -q
25772  25588 sh -c test -f '/usr/bin/python-config' -a -x '/usr/bin/python-config' && echo y
25773  25588 /bin/sh -c '/usr/bin/python-config' --ldflags 2>/dev/null
25774  25773 /usr/bin/python-config --ldflags
25775  25588 /bin/sh -c '/usr/bin/python-config' --cflags 2>/dev/null
25776  25775 /usr/bin/python-config --cflags
25777  25588 /bin/sh -c touch PERF-FEATURES; cat PERF-FEATURES
25778  25777 touch PERF-FEATURES
25779  25777 cat PERF-FEATURES
25780  25588 printf ...%30s: [ \033[32mon\033[m ] dwarf
25781  25588 printf ...%30s: [ \033[31mOFF\033[m ] dwarf
25782  25588 printf ...%30s: [ \033[32mon\033[m ] glibc
25783  25588 printf ...%30s: [ \033[31mOFF\033[m ] glibc
25784  25588 printf ...%30s: [ \033[32mon\033[m ] gtk2
[...]
```

Wow. Application startups can also run a surprising number of processes. execsnoop can help you identify areas for performance improvement: excessive sh/grep/sed/awk invocations, that can often be rewritten to use more advanced shell or awk features.

You can also use execsnoop to catch unexpected behavior. Running it with -t for timestamps:

```
# ./execsnoop -t
Tracing exec(s). Ctrl-C to end.
TIMES      PID  PPID  ARGS
[...]
```

3799932.757407	29403	29390	hostname
3799932.763324	29405	29404	cat /sys/class/net/eth0/address
3799932.768432	29407	29406	grep -l ^[[:space:]]*EC2SYNC=no\([[:space:]]#\ \$\) /etc/sysconfig/
3799932.778556	29412	29411	curl -s -f http://169.254.169.254/latest/meta-data/network/interfa
3799932.796555	29415	29413	grep inet .* secondary eth0
3799932.798570	29416	29413	awk {print \$2}
3799932.801021	29414	29413	/sbin/ip addr list dev eth0 secondary
3799932.802559	29417	29413	cut -d/ -f1
3800159.574170	29419	29418	/bin/sh -c /usr/lib64/sa/sa1 1 1

Why is my system running curl on that address? Hm.

## Options

[execsnoop](#) options are summarized by the USAGE message (there's also a [man page](#) and [examples file](#)):

```
# ./execsnoop -h
USAGE: execsnoop [-hrt] [-a argc] [-d secs] [name]
      -d seconds      # trace duration, and use buffers
      -a argc         # max args to show (default 8)
      -r              # include re-execs
      -t              # include time (seconds)
      -h              # this usage message
      name            # process name to match (REs allowed)

eg,
execsnoop             # watch exec(s) live (unbuffered)
execsnoop -d 1        # trace 1 sec (buffered)
execsnoop grep        # trace process names containing grep
execsnoop 'log$'      # filenames ending in "log"
```

execsnoop traces events as they happen, unless the -d option is used, which uses in-kernel buffering.

## What, Why, and How

This is another ftrace-based hack for my [perf-tools](#) collection. These are designed to work with fewest dependencies (including no kernel debuginfo, if possible), and on older Linux kernel versions, particularly my Linux 3.2 systems. I expect to rewrite them when new tracing features are added to Linux in the future.

This turned out to be difficult, and a number of times it I thought it might be impossible in this environment.

My first attempt, [execsnoop-proc](#), traced sched:sched\_process\_exec with process arguments from /proc/PID/cmdline. This worked on many systems, but not all, as I don't think /proc could always be read quickly enough for some processes. The sched:sched\_process\_exec tracepoint was missing on some systems as well.

The version I'm now using dynamically traces either stub\_execve() or do\_execve(), and walks the %si register as an array of strings using an unrolled loop. This is an enormous hack that I can hardly believe works, but it does work on all the systems I need it to.

Here's the essence of what I'm doing, using the perf(1) command ([perf events](#)):

```
# perf probe --add 'do_execve +0(+0(%si)):string +0(+8(%si)):string +0(+16(%si)):string +0(+24(%si)):string'
# perf record --no-buffering -e probe:do_execve -a -o - | PAGER="cat -v" stdbuf -oL perf script -
:10007 10007 [000] 557516.214765: probe:do_execve: (ffffffff811cccb0) arg1="ls" arg2="--color=auto"
:10008 10008 [000] 557516.219168: probe:do_execve: (ffffffff811cccb0) arg1="sleep" arg2="1" arg3=""
# perf probe --del do_execve
```

In that case, I only included three arguments, but you can see how the unrolled loop works. The first one-liner is for a specific kernel version and platform, and may not work for you for many reasons, without first adjusting it to match what you have.

If you want to analyze at this level, you might find my [kprobe](#) tool easier to work with. Compare the above to:

```
# ./kprobe 'p:do_execve +0(+0(%si)):string +0(+8(%si)):string +0(+16(%si)):string +0(+24(%si)):string'
Tracing kprobe do_execve. Ctrl-C to end.
kprobe-12484 [000] d... 7149163.889695: do_execve: (do_execve+0x0/0x20) arg1="cat" arg2="trace"
bash-12499 [001] d... 7149164.593106: do_execve: (do_execve+0x0/0x20) arg1="ls" arg2="--color=auto"
bash-12500 [001] d... 7149164.597399: do_execve: (do_execve+0x0/0x20) arg1="sleep" arg2="1" arg3=""
^C
Ending tracing...
```

kprobe(8) is nice. It automatically adds and removes kprobes, and has options for showing stacks (-s) and column headers (-H).

Were I to use SystemTap, I'd do something like this:

```
# stap -ve 'probe process.begin { printf("%6d %6d %s\n", pid(), ppid(), cmdline_str()); }'
```

(When I run this, for some reason it begins by listing all the current processes, followed by the new ones. I don't know if it's a feature or a bug.)

The future of Linux should support a tracer that does a one-liner like this. In the meantime, I can use this hacked version of execsnoop to solve performance issues on older systems.

## Conclusion

Short-lived processes can cause performance problems, and aren't visible from interval-sampling tools like top(1). In this post I described a Linux port of my popular execsnoop tool, which I've used for many years to identify and study short-lived processes, and solve the performance problems they can cause.

This is also another proof of concept for older Linux kernels, like [opensnoop](#) and [iosnoop](#), using the existing ftrace and kprobes tracing frameworks. There are other ways to do this: you might have kernel debuginfo and SystemTap available, or auditing enabled. This implementation is handy for my basic Linux 3.2 cloud instances, that don't have kernel debuginfo installed.

Warnings apply: execsnoop and these one-liners use dynamic tracing on Linux, which has had kernel panic bugs in the past, so know what you are doing, test first, and use at your own risk.

---

*You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).*

---

Copyright 2017 Brendan Gregg.

[About this blog](#)

[Perf Metrics](#)

[Off-CPU Analysis](#)

[Active Bench.](#)

[Flame Graphs](#)

[Heat Maps](#)

[Frequency Trails](#)

[Colony Graphs](#)

[perf Examples](#)

[eBPF Tools](#)

[DTrace Tools](#)

[DTraceToolkit](#)

[DtkshDemos](#)

[Guessing Game](#)

[Specials](#)

[Books](#)

[Other Sites](#)