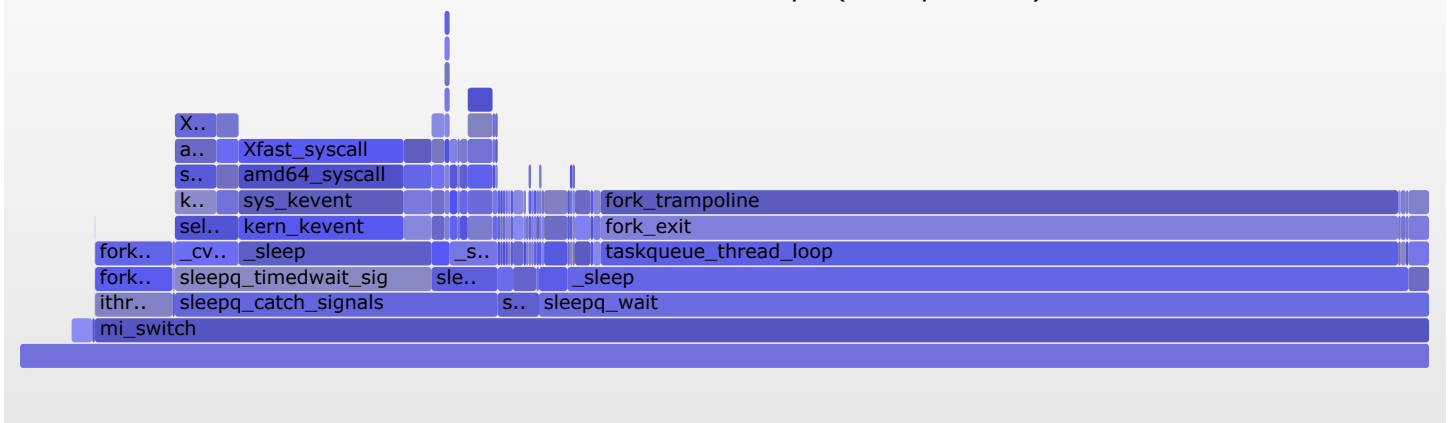


## FreeBSD Off-CPU Flame Graphs

12 Mar 2015

While I was giving a talk on FreeBSD flame graphs, a new technique was suggested by the audience for gathering off-CPU stacks, which is something I'd like to see improved and included in all operating systems. It involved using `procstat -ka` to gather sleeping thread stacks, from which I created an off-CPU flame graph:

## FreeBSD Off-CPU Flame Graph (from procstat)



This visualization lets us browse the sleeping thread call stacks to look for anything unusual that may be a problem. Most of these will be normal, as threads usually spend most of their time sleeping while waiting for work. Some may not be normal, and can be the source of latency and performance issues.

## Background

CPU profiling can help solve all performance issues when threads are on-CPU, but what about issues when threads are blocked and sleeping off-CPU? Time may be spent waiting for storage I/O, network I/O, lock contention, run queue latency, etc. Collectively, I call this *off-CPU time*, which can be profiled and examined to help solve these issues.

I recently showed how [Linux perf events](#) can measure off-CPU time and create flame graphs. This approach captures all kernel scheduler events involved in sleeping and switching threads, to a trace file, and then uses user-level programs to stitch together events, calculating off-CPU time by call stack. While it works, tracing all scheduler events to a file can incur crazy-high overhead, and I'd be very nervous about using this on a busy production system.

An improvement is to calculate off-CPU time in kernel context, and to pass only aggregated call stacks with total times to user-level. This needs a programmable tracer, like SystemTap on Linux, or DTrace on FreeBSD. [Slide 48](#) of my talk showed the DTrace commands, but this still comes with a high overhead warning, even though it's much better than the `perf_events` approach.

Off-CPU tracing costs overhead that is relative to the event rate. A system doing over a million IOPS can have over a million scheduler events per second. Compare this to CPU profiling, where we wake up at a fixed rate (eg, 99 Hertz) on all CPUs and gather a stack sample. On an 8 CPU system, CPU profiling can mean about 800 stack samples per second, no matter how busy the server is. Not the millions per second that off-CPU tracing can reach.

Can we profile off-CPU time in the same way as CPU profiling? Instead of tracing every scheduler event, can we wake up at a timed interval and gather all the *sleeping* stacks, not the on-CPU ones? This would mean the overhead was relative to the thread count instead of the scheduler event rate. Eg, to sample these at, say, 9 Hertz, on a system with 500 threads, would mean taking about 4,500 stack samples per second.

## FreeBSD procstat -ka

During my talk, Adrian Chadd suggested I use `procstat -ka` to dump sleeping stacks (thanks!):


```
# procstat -ka
PID      TID COMM          TDNAME          KSTACK
0 100000 kernel        swapper         mi_switch sleepq_timedwait _sleep swapper btext
0 100023 kernel        firmware taskq  mi_switch sleepq_wait _sleep taskqueue_thread_loo
0 100025 kernel        kqueue taskq    mi_switch sleepq_wait _sleep taskqueue_thread_loo
0 100026 kernel        ffs trim taskq  mi_switch sleepq_wait _sleep taskqueue_thread_loo
0 100027 kernel        acpi_task_0     mi_switch sleepq_wait msleep_spin_sbt taskqueue_th
0 100028 kernel        acpi_task_1     mi_switch sleepq_wait msleep_spin_sbt taskqueue_th
0 100029 kernel        acpi_task_2     mi_switch sleepq_wait msleep_spin_sbt taskqueue_th
0 100033 kernel        aiobio taskq    mi_switch sleepq_wait _sleep taskqueue_thread_loo
0 100035 kernel        thread taskq    mi_switch sleepq_wait _sleep taskqueue_thread_loo
0 100037 kernel        mps0 taskq      mi_switch sleepq_wait _sleep taskqueue_thread_loo
0 100049 kernel        mps1 taskq      mi_switch sleepq_wait _sleep taskqueue_thread_loo
[...]
```

Wow, so FreeBSD *already has this*. I'd really like to do this on Linux, too (gdb macro, at least?).

The KSTACK column shows a leaf to origin call stack. It just needs a little text processing for flamegraph.pl. For example, collecting ten of these and generating an off-CPU flame graph (in Bourne shell):

```
# git clone --depth 1 https://github.com/brendangregg/FlameGraph
# cd FlameGraph
# for i in `seq 1 10`; do procstat -ka >> out.offstacks01; sleep 1; done
# sed 's/^\.{47}\///;s/ /;/g;s/;/$/ 1/' out.offstacks01 | \
./flamegraph.pl --color=io --title="Off-CPU" > offstacks01.svg
```

That generates the off-CPU flame graph I demonstrated in the talk, and is also at the top of this page ([SVG](#)):



Click to zoom.

## By Thread


Since procstat has the process thread names as the third and fourth fields, we can include them too. This gets a bit trickier as the thread name can contain spaces. Switching to awk:

```
# awk '{ comm = $3; thread = substr($0, 30, 16); gsub(/ *$/, "", thread);
stack = substr($0, 48); gsub(/ /, ";", stack); gsub(/;$/, "", stack);
printf "%s;%s;%s\n", comm, thread, stack }' out.offstacks01 | \
./flamegraph.pl --color=io --title="Off-CPU" > offstacks01byt.svg
```

(This awk is getting a bit long, and I should put it in a stackcollapse-procstat.awk file, like I have for other profiler output.)

Example output, now including process and thread names ([SVG](#)):

## FreeBSD Off-CPU Flame Graph



This has more detail, as we're now splitting time for each thread name, creating many thin towers. The wider towers mean pools of threads. There are 64 threads in ZFS metaslab\_group\_t (truncated), and 65 in nginx.

The per-thread name detail might be overdoing it in many cases. You can adjust the awk to only include the process name if desired.

## Improving procstat


There is room for improvement with procstat. First, its stacks are kernel only. It would be helpful for it to include user-level frames, to fully explain blocked events.

Second, the overhead is higher than I would have hoped:

```
# procstat -ka | wc -l
546
# time procstat -ka > /dev/null
1.27 real    0.00 user    1.27 sys
```

That's 1.27 seconds of system CPU time to walk 546 threads. We won't be able to take even one sample of all threads per second, let alone taking it to 9 Hertz (like I suggested earlier) or higher.

Here's a CPU flame graph of procstat's system CPU time:



Ok, so we're in link\_elf\_search\_symbol() (that's so obvious I didn't need to visualize this as a flame graph).

From sys/kern/link\_elf.c:

```
static int
link_elf_search_symbol(linker_file_t lf, caddr_t value,
    c_linker_sym_t *sym, long *diffp)
{
    [...]
    for (i = 0, es = ef->dbsymtab; i < ef->dbsymcnt; i++, es++) {
        if (es->st_name == 0)
            continue;
        st_value = es->st_value + (uintptr_t) (void *) ef->address;
        if (off >= st_value) {
            if (off - st_value < diff) {
                diff = off - st_value;
                best = es;
                if (diff == 0)
                    break;
            } else if (off - st_value == diff) {
                best = es;
            }
        }
    }
}
```

This is not only a sequential search of a list for a symbol, but it also checks the entire list unless a direct match is found (offset 0), which usually won't be the case. Usually, it'll return the closest match. This is also called by another loop in `linker_debug_search_symbol()` (not visible in the flame graph, probably due to compiler optimizations), which iterates over all linker symbol files. In this case, it's searching all the kernel modules, the largest of which is `zfs`, which has 4,532 symbols.

So for every frame in every stack trace, we usually compare the program counter to every known kernel symbol, sequentially. This can be improved: eg, using a btree lookup, or at least bailing out early when a decent match is found. Another approach could be to build an off-CPU profiler in the kernel, which sampled and then aggregated stacks before doing symbol lookup. Many approaches are possible. I'd suspect there's never been a reason to optimize this further until now.

## Conclusion

FreeBSD has an interesting way to list blocked thread stacks using `procstat -ka`, which can be used to generate off-CPU flame graphs. In this post I described how it worked, and some room for improvement with this approach. It would be great to see similar facilities in other operating systems, so that off-CPU flame graphs can be generated in a cheaper fashion, for solving issues of blocked time.

For more about flame graphs, see my previous post about my [FreeBSD flame graphs talk](#).

Comments for this thread are now closed ×

0 Comments

Brendan Gregg's Blog

1

Login

Recommend

2

Tweet

Share

Sort by Best

This discussion has been closed.

Subscribe
 Add Disqus to your site
 Add Disqus
 Add
 Disqus' Privacy Policy
 Privacy Policy
 Privacy Policy
 Privacy Policy

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if disqus add advertisements).

[Homepage](#)

[Blog](#)

[Full Site Map](#)

[Sys Perf book](#)

[Linux Perf](#)

[Perf Methods](#)

[USE Method](#)

[TSA Method](#)

[Off-CPU Analysis](#)

[Active Bench.](#)

[Flame Graphs](#)

[Heat Maps](#)

[Frequency Trails](#)

[Colony Graphs](#)

[perf Examples](#)

[eBPF Tools](#)

[DTrace Tools](#)

[DTraceToolkit](#)

[DtkshDemos](#)

[Guessing Game](#)

[Specials](#)

[Books](#)

[Other Sites](#)