

ftrace: The Hidden Light Switch

30 Aug 2014

If you've been following my previous posts about Linux ftrace, you'll want to read my recent [lwn.net](#) article: [Ftrace: The Hidden Light Switch](#). This showed how I used ftrace at Netflix to confirm whether a kernel tunable change took affect immediately, and how it was set. It was also a tour of some ftrace capabilities.

ftrace is particularly useful for us since it is already available in all our Linux systems, and has been for years. The biggest problem is that it's not very well known. I've hopefully helped raise awareness with this article, and I also included ftrace in my [Linux Performance Tools](#) talk at LinuxCon.

I mentioned the following commands, but didn't include screenshots. Here I'm checking if the deadline or noop I/O schedulers are active, based on counting the function calls:

```
# ./funccount -i 1 'deadline*'
Tracing "deadline*"... Ctrl-C to end.

FUNC                                COUNT
deadline_add_request               178
deadline_merge                     178
deadline_remove_request.isra.4     178
deadline_dispatch_requests          712

FUNC                                COUNT
deadline_add_request               698
deadline_merge                     698
deadline_remove_request.isra.4     698
deadline_dispatch_requests          2792
^C
FUNC                                COUNT
deadline_add_request               316
deadline_merge                     316
deadline_remove_request.isra.4     316
deadline_dispatch_requests          1264

Ending tracing...
# ./funccount -i 1 'noop*'
Tracing "noop*"... Ctrl-C to end.

FUNC                                COUNT

FUNC                                COUNT
^C
FUNC                                COUNT

Ending tracing...
```

So, deadline is active right now...

These and the other ftrace capabilities have been invaluable so far, which we've been using via my [perf-tools](#) wrappers. If you want to do more with ftrace, see [trace-cmd](#) by Steven Rostedt, which is a more powerful multi-tool.

ftrace can't yet do everything that I want from a kernel tracer. For example, I can't yet do custom in-kernel aggregations. But it can do a lot more than I previously thought possible.

Thanks to [Deirdré Straughan](#) and Jonathan Corbet for editing help with the article. For more about ftrace, see the resources I linked to at the end (many also on [lwn.net](#)), and my previous posts on [ftrace](#).

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).

Copyright 2017 Brendan Gregg.

[About this blog](#)

[Perf Methods](#)

[USE Method](#)

[TSA Method](#)

[Off-CPU Analysis](#)

[Active Bench.](#)

[Flame Graphs](#)

[Heat Maps](#)

[Frequency Trails](#)

[Colony Graphs](#)

[perf Examples](#)

[eBPF Tools](#)

[DTrace Tools](#)

[DTraceToolkit](#)

[DtkshDemos](#)

[Guessing Game](#)

[Specials](#)

[Books](#)

[Other Sites](#)