

Linux bcc Tracing Security Capabilities

01 Oct 2016

Which Linux security capabilities are your applications using? I recently developed a new tool, capable, to print out capability checks live:

```
# capable
TIME      UID      PID      COMM      CAP      NAME      AUDIT
22:11:23  114      2676     snmpd      12      CAP_NET_ADMIN  1
22:11:23   0      6990     run       24      CAP_SYS_RESOURCE  1
22:11:23   0      7003     chmod      3      CAP_FOWNER  1
22:11:23   0      7003     chmod      4      CAP_FSETID  1
22:11:23   0      7005     chmod      4      CAP_FSETID  1
22:11:23   0      7005     chmod      4      CAP_FSETID  1
22:11:23   0      7006     chown      4      CAP_FSETID  1
22:11:23   0      7006     chown      4      CAP_FSETID  1
22:11:23   0      6990     setuidgid   6      CAP_SETGID  1
22:11:23   0      6990     setuidgid   6      CAP_SETGID  1
22:11:23   0      6990     setuidgid   7      CAP_SETUID  1
22:11:24   0      7013     run       24      CAP_SYS_RESOURCE  1
22:11:24   0      7026     chmod      3      CAP_FOWNER  1
[...]
```

capable uses [bcc](#), a front-end and a collection of tools that use new Linux enhanced BPF tracing capabilities. capable works by using BPF with kprobes to dynamically trace the kernel `cap_capable()` function, and then uses a table to map the capability index to the name seen in the output. Here's the [source code](#): it's pretty straightforward.

I wrote it as a colleague, Michael Wardrop, asked what security capabilities our applications were actually using. Given a list, we could use `setcap(8)` (or other software) to improve the security of applications by only allowing the necessary capabilities.

Non-audit Checks

The `cap_capable()` function has an audit argument, which directs whether the capability check should write an audit message or not, if that's configured. By default, I only print capability checks where this is true, but capable can also trace all checks with the `-v` option:

```
# capable -h
usage: capable [-h] [-v] [-p PID]

Trace security capability checks

optional arguments:
  -h, --help            show this help message and exit
  -v, --verbose          include non-audit checks
  -p PID, --pid PID     trace this PID only

examples:
  ./capable             # trace capability checks
  ./capable -v          # verbose: include non-audit checks
  ./capable -p 181      # only trace PID 181
```

Here's some non-audit events:

```
# capable -v
TIME      UID      PID      COMM      CAP      NAME      AUDIT
20:53:45  60004    22061    lsb_release  21      CAP_SYS_ADMIN  0
20:53:45  60004    22061    lsb_release  21      CAP_SYS_ADMIN  0
20:53:45  60004    22061    lsb_release  21      CAP_SYS_ADMIN  0
20:53:45  60004    22061    lsb_release  21      CAP_SYS_ADMIN  0
20:53:45  60004    22061    lsb_release  21      CAP_SYS_ADMIN  0
20:53:45  60004    22061    lsb_release  21      CAP_SYS_ADMIN  0
[...]
```

What are all those?

I'll start by showing the `cap_capable()` function prototype, from `security/commoncap.c`:

```
int cap_capable(const struct cred *cred, struct user_namespace *targ_ns,
                int cap, int audit)
```

Now I can use `bcc`'s trace program to inspect these calls (bear with me), given that `cap` will be `arg3`, and `audit` `arg4` (from the above prototype):

```
# trace 'cap_capable "cap: %d, audit: %d", arg3, arg4'
TIME    PID    COMM      FUNC
20:56:18 25535  lsb_release  cap_capable  cap: 21, audit: 0
20:56:18 25535  lsb_release  cap_capable  cap: 21, audit: 0
20:56:18 25535  lsb_release  cap_capable  cap: 21, audit: 0
20:56:18 25535  lsb_release  cap_capable  cap: 21, audit: 0
20:56:18 25535  lsb_release  cap_capable  cap: 21, audit: 0
[...]
```

That one-liner is pretty similar to my capable program, except it lacks the "NAME" column with human readable translations.

I'm really doing this so I can add the (newly added) -K and -U options, which print kernel and user-level stack traces. I'll just use -K:

```
# trace -K 'cap_capable "cap: %d, audit: %d", arg3, arg4'
TIME    PID    COMM      FUNC
[...]
```

20:59:58 30607 lsb_release cap_capable cap: 21, audit: 0

Kernel Stack Trace:

```
ffffff813659d1 cap_capable
ffffff813684bb security_vm_enough_memory_mm
ffffff811deda6 expand_downwards
ffffff811def64 expand_stack
ffffff81234321 setup_arg_pages
ffffff8128c10b load_elf_binary
ffffff81234cee search_binary_handler
ffffff8128b7ff load_script
ffffff81234cee search_binary_handler
ffffff8123635a do_execveat_common.isra.35
ffffff812367da sys_execve
ffffff81003bae do_syscall_64
ffffff81861ca5 return_from_SYSCALL_64
```

20:59:58 30607 lsb_release cap_capable cap: 21, audit: 0

Kernel Stack Trace:

```
ffffff813659d1 cap_capable
ffffff813684bb security_vm_enough_memory_mm
ffffff811ddf623 mmap_region
ffffff811dff4b do_mmap
ffffff811c122a vm_mmap_pgoff
ffffff811c1295 vm_mmap
ffffff8128bb93 elf_map
ffffff8128c271 load_elf_binary
ffffff81234cee search_binary_handler
ffffff8128b7ff load_script
ffffff81234cee search_binary_handler
ffffff8123635a do_execveat_common.isra.35
ffffff812367da sys_execve
ffffff81003bae do_syscall_64
ffffff81861ca5 return_from_SYSCALL_64
```

[...]

Awesome. So these are coming from `security_vm_enough_memory_mm()`. By reading the source, I see it's used to reserve some memory for root. It's not a hard failure if the capability is missing. And it's not really a security check, hence why it disabled audit.

I should add a -K option to the capable tool, so it can print stack traces too.

Older Kernels

To use capable, you'll need a 4.4 kernel. To use the -K option, 4.6.

Here's a version using my older [perf-tools](#) collection, which uses `ftrace` and should work on much older kernels including the 3.x series:

```
# ./perf-tools/bin/kprobe -s 'p:cap_capable cap=%dx audit=%cx' 'audit != 0'
Tracing kprobe cap_capable. Ctrl-C to end.
      run-4440 [003] d... 6417394.367486: cap_capable: (cap_capable+0x0/0x70) cap=0x18 aud
      run-4440 [003] d... 6417394.367492:
=> ns_capable_common
=> capable
=> do_prlimit
=> Sys_setrlimit
=> entry_SYSCALL_64_fastpath
      chmod-4453 [006] d... 6417394.399020: cap_capable: (cap_capable+0x0/0x70) cap=0x3 aud
      chmod-4453 [006] d... 6417394.399027:
=> ns_capable_common
=> ns_capable
=> inode_owner_or_capable
=> inode_change_ok
=> xfs_setattr_nonsize
=> xfs_vn_setattr
=> notify_change
=> chmod_common
=> Sys_fchmodat
=> entry_SYSCALL_64_fastpath
      chmod-4453 [006] d... 6417394.399035: cap_capable: (cap_capable+0x0/0x70) cap=0x4 aud
      chmod-4453 [006] d... 6417394.399037:
=> ns_capable_common
=> capable_wrt_inode_uidgid
=> inode_change_ok
=> xfs_setattr_nonsize
=> xfs_vn_setattr
=> notify_change
=> chmod_common
=> Sys_fchmodat
=> entry_SYSCALL_64_fastpath
      chmod-4455 [007] d... 6417394.402524: cap_capable: (cap_capable+0x0/0x70) cap=0x4 aud
      chmod-4455 [007] d... 6417394.402529:
=> ns_capable_common
=> capable_wrt_inode_uidgid
=> inode_change_ok
=> xfs_setattr_nonsize
=> xfs_vn_setattr
=> notify_change
=> chmod_common
=> Sys_fchmodat
=> entry_SYSCALL_64_fastpath
[...]
```

It's a one-liner using my kprobe tool. It's also (currently) a bit harder to use: I need to know which registers those arguments will be in: the example above is for x86_64 only.

That's all for now. Happy hacking.

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).