

Linux bcc/BPF tcplife: TCP Lifespans

30 Nov 2016

"i really wish i had a command line tool that would give me stats on TCP connection lengths on a given port"

```
# ./tcplife -D 80
PID  COMM      LADDR          LPORT  RADDR          RPORT  TX_KB  RX_KB  MS
27448 curl      100.66.11.247  54146  54.154.224.174  80      0      1  263.85
27450 curl      100.66.11.247  20618  54.154.164.22   80      0      1  243.62
27452 curl      100.66.11.247  11480  54.154.43.103   80      0      1  231.16
27454 curl      100.66.11.247  31382  54.154.15.7     80      0      1  249.95
27456 curl      100.66.11.247  33416  52.210.59.223   80      0      1  545.72
27458 curl      100.66.11.247  16406  52.30.140.35    80      0      1  222.29
27460 curl      100.66.11.247  11634  52.30.133.135   80      0      1  217.52
27462 curl      100.66.11.247  25660  52.30.126.182   80      0      1  250.81
[...]
```

That's tracing destination port 80, you can also trace local port (-L), or trace all ports (default behavior).

The quote and good idea is from Julia Evans on [twitter](#), and I added it as a tool to the Linux BPF-based [bcc](#) open source collection.

The output of tcplife, short for TCP lifespan, shows not just the duration (MS == milliseconds) but also throughput statistics: TX_KB for Kbytes transmitted, and RX_KB for Kbytes received. It should be useful for performance and security analysis, and network debugging.

I am NOT tracing packets

The overheads can cost too much to examine every packet, especially on servers that process millions of packets per second. To do this I'm not using tcpdump, libpcap, tethereal, or any network sniffer.

So how'd I do it? There were four challenges:

1. Measuring lifespans

The current version of tcplife uses kernel dynamic tracing (kprobes) of `tcp_set_state()`, and looks for the duration from an early state (eg, `TCP_ESTABLISHED`) to `TCP_CLOSE`. State changes have a much lower frequency than packets, so this approach greatly reduces overhead. But it ends up trickier than it sounds, since it's tracing the Linux implementation of TCP, which isn't guaranteed to use `tcp_set_state()` for every state transition. But it works well enough for now, and we can add a stable tracepoint for TCP state transitions to future kernels so that this will always work. (Or, if that proves untenable, tracepoints for the creation and destruction of TCP sessions or sockets.)

2. Fetching addresses and ports

`tcp_set_state()` has `struct sock *sk` as an argument, and we can dig the details from that.

3. Fetching throughput statistics

Those TX_KB and RX_KB columns. This was only possible in the last couple of years thanks to the RFC-4898 additions to `struct tcp_info` in the Linux kernel: `tcpi_bytes_acked` and `tcpi_bytes_received`. I discussed these in my [tcptop](#) blog post.

4. Showing task context

It's useful to show the PID and COMM (process name) with these connections, but TCP state changes aren't guaranteed to happen in the correct task context, so we can't just fetch the currently running task information. Nor is there a cached PID and comm in `struct sock` to print out. So I'm caching the task context on TCP state changes where it's usually valid, by virtue of implementation. It works, but is another area where we can do better, and can be addressed if and when we add stable TCP tracepoints.

That's how tcplife works right now. In the future it may change and improve, especially if more TCP tracepoints become available. I'm also not the first to try this kind of TCP event tracing: Facebook already have their own BPF TCP event tool, although I think their implementation is a little different.

Here's the full USAGE message for tcplife:

```

# ./tcplife -h
usage: tcplife [-h] [-T] [-t] [-w] [-s] [-p PID] [-L LOCALPORT]
              [-D REMOTEPORT]

Trace the lifespan of TCP sessions and summarize

optional arguments:
  -h, --help            show this help message and exit
  -T, --time            include time column on output (HH:MM:SS)
  -t, --timestamp       include timestamp on output (seconds)
  -w, --wide            wide column output (fits IPv6 addresses)
  -s, --csv             comma separated values output
  -p PID, --pid PID     trace this PID only
  -L LOCALPORT, --localport LOCALPORT
                        comma-separated list of local ports to trace.
  -D REMOTEPORT, --remoteport REMOTEPORT
                        comma-separated list of remote ports to trace.

examples:
  ./tcplife             # trace all TCP connect()s
  ./tcplife -t          # include time column (HH:MM:SS)
  ./tcplife -w          # wider columns (fit IPv6)
  ./tcplife -stT        # csv output, with times & timestamps
  ./tcplife -p 181       # only trace PID 181
  ./tcplife -L 80        # only trace local port 80
  ./tcplife -L 80,81     # only trace local ports 80 and 81
  ./tcplife -D 80        # only trace remote port 80

```

And there is more [example output](#) in bcc, along with a [man page](#). For more about bcc and BPF (eBPF), see the collection of documents on my [Linux performance page](#).

The only catch is that tcplife does need newer kernels (say, 4.4).

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).