

Linux iosnoop Latency Heat Maps

23 Jul 2014

Odd patterns of I/O latency can be hidden by line graphs and summary statistics, and revealed by histograms and heat maps. In my previous post I showed my Linux [iosnoop](#) tool, which can trace block device I/O along with timestamps and latency. This information can be visualized, revealing any odd patterns.

As an example, I'll make a latency heat map from iosnoop output using my [trace2heatmap.pl](#) program (which is also on [github](#)). Here's the command I used, which captures both start (-s) and end (-t) timestamps:

```
# ./iosnoop -ts 90 > out.iosnoop
# more out.iosnoop
Tracing block I/O for 90 seconds (buffered)...
STARTs      ENDs      COMM      PID    TYPE  DEV      BLOCK      BYTES    LATms
6743904.592147 6743904.592316 java      9823    R     202,32    17266904    8192     0.17
6743904.594729 6743904.594907 java      9823    R     202,16    23030152    8192     0.18
6743904.597172 6743904.597402 java      9823    R     202,32    1405848     8192     0.23
6743904.598571 6743904.598745 java      9823    R     202,32    25259784    8192     0.17
[...]
```


I won't use the start timestamp (STARTs), but having it at high resolution may be useful for later study.

Now converting this into a heatmap:

```
# git clone https://github.com/brendangregg/HeatMap
# cd HeatMap
# cat ../out.iosnoop | awk '$1 ~ /^[0-9]/ { print $2, $9 }' | ./trace2heatmap.pl \
  --unitslatency=ms --unitstime=s --maxcol=90 --maxlat=100 --grid \
  --title="Block I/O Latency Heat Map" > biolatenyheatmap.svg
```

You can tune the various options as desired. I truncated it to 100 milliseconds (--maxlat), and added a title. This makes the following (mouse-over for details, or try the direct [SVG](#)):

Block I/O Latency Heat Map



Great! I use these to examine latency outliers, as well as the distribution for the bulk of the I/O. Time is on the x-axis, and I/O latency on the y-axis. The number of I/O at each time and latency range is shown by the darkness of each block: darker for more.

In this case the bulk of the I/O is very fast, between 0 and 2 milliseconds (shown as the dark red line at the bottom). There are also clouds of high latency I/O, about every 9 seconds, which are around 70 milliseconds, creating a multimodal distribution. Their presence would be difficult to see from average latency alone.

It turns out that these are due to a single disk in particular, which I can filter using awk:

```
# cat ../out.iosnoop | awk '$6 == "202,1" { print $2, $9 }' | ...
```

I also reduced the queue size to this slow disk using:

```
# echo 4 > /sys/block/xvda1/queue/nr_requests
```

This type of tuning was suggested in my previous blog's comments, as a possible relief for I/O latency outliers caused by reads queueing behind a large batch of writes. This reduced the severity of the I/O latency outliers a little, and the I/O clouds a lot. (This tuning will also hurt performance for that one disk, so don't copy it without understanding what it does.)

See the [before \(128 queue length\)](#) and [after \(4 queue length\)](#) latency heat maps for that disk.

Have disk I/O issues? Aren't using latency heat maps? You should! See my [heat maps](#) page for more details.

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).