

## Linux bcc/eBPF tcpdrop

31 May 2018

While debugging a production issue of kernel-based TCP packet drops, I remembered seeing a new function added in Linux 4.7 by Eric Dumazet (Google) called `tcp_drop()`, which I can trace using kprobes and bcc/eBPF. This lets me fetch extra context to explain why these drops are happening. Eg:

```
# tcpdrop
TIME      PID      IP SADDR:SPORT      > DADDR:DPORT      STATE (FLAGS)
05:46:07 82093   4   10.74.40.245:50010 > 10.74.40.245:58484 ESTABLISHED (ACK)
tcp_drop+0x1
tcp_rcv_established+0x1d5
tcp_v4_do_rcv+0x141
tcp_v4_rcv+0x9b8
ip_local_deliver_finish+0x9b
ip_local_deliver+0x6f
ip_rcv_finish+0x124
ip_rcv+0x291
__netif_receive_skb_core+0x554
__netif_receive_skb+0x18
process_backlog+0xba
net_rx_action+0x265
__softirqentry_text_start+0xf2
irq_exit+0xb6
xen_evtchn_do_upcall+0x30
xen_hvm_callback_vector+0x1af

05:46:07 85153   4   10.74.40.245:50010 > 10.74.40.245:58446 ESTABLISHED (ACK)
tcp_drop+0x1
tcp_rcv_established+0x1d5
tcp_v4_do_rcv+0x141
tcp_v4_rcv+0x9b8
ip_local_deliver_finish+0x9b
ip_local_deliver+0x6f
ip_rcv_finish+0x124
ip_rcv+0x291
__netif_receive_skb_core+0x554
__netif_receive_skb+0x18
process_backlog+0xba
net_rx_action+0x265
__softirqentry_text_start+0xf2
irq_exit+0xb6
xen_evtchn_do_upcall+0x30
xen_hvm_callback_vector+0x1af

[...]
```

This is [tcpdrop](#), a new tool I've written for the open source [bcc](#) project. It shows source and destination packet details, as well as TCP session state (from the kernel), TCP flags (from the packet TCP header), and the kernel stack trace that led to this drop. That stack trace helps answer the why (you'll need to browse the code behind those functions to make sense of it). This is also information that's not on the wire, so you can never see this using packet sniffers (libpcap, tcpdump, etc).

I can't help but highlight this small but significant change from Eric's patch ([tcp: increment sk\\_drops for dropped rx packets](#)):

```
@@ -6054,7 +6061,7 @@ int tcp_rcv_state_process(struct sock *sk, struct sk_buff *skb)
    if (!queued) {
        discard:
-        __kfree_skb(skb);
+        tcp_drop(sk, skb);
    }
    return 0;
```

`__kfree_skb()` is called from many paths to free socket buffers, including routine codepaths. Tracing it was too noisy: you'd have your packet drop code paths lost among many normal ones. But with the new `tcp_drop()` function, I can trace just the TCP drops. I've already suggested some enhancements to Eric today at netconf18, such as adding a "reason" argument somewhere that I can trace for a more human description of why the packet was dropped. Maybe `tcp_drop()` should become a tracepoint too.

Here's some more code worth mentioning, this time some eBPF/C from my `tcpdrop` tool:

```
[...]
// pull in details from the packet headers and the sock struct
u16 family = sk->__sk_common.skc_family;
char state = sk->__sk_common.skc_state;
u16 sport = 0, dport = 0;
u8 tcpflags = 0;
struct tcphdr *tcp = skb_to_tcphdr(skb);
struct iphdr *ip = skb_to_iphdr(skb);
bpf_probe_read(&sport, sizeof(sport), &tcp->source);
bpf_probe_read(&dport, sizeof(dport), &tcp->dest);
bpf_probe_read(&tcpflags, sizeof(tcpflags), &tcp_flag_byte(tcp));
sport = ntohs(sport);
dport = ntohs(dport);

if (family == AF_INET) {
    struct ipv4_data t data4 = {.pid = pid, .ip = 4};
    bpf_probe_read(&data4.saddr, sizeof(u32), &ip->saddr);
    bpf_probe_read(&data4.daddr, sizeof(u32), &ip->daddr);
    data4.dport = dport;
    data4.sport = sport;
    data4.state = state;
    data4.tcpflags = tcpflags;
    data4.stack_id = stack_traces.get_stackid(ctx, 0);
    ipv4_events.perf_submit(ctx, &data4, sizeof(data4));
}
[...]
```

My prior tcp tools in bcc have made do with struct sock members alone (eg, [tcp life](#)). But this time I needed packet info to see TCP flags, and the direction of the packet. So it's the first time I've accessed TCP and IP headers in eBPF. I added `skb_to_tcphdr()` and `skb_to_iphdr()` in `tcpdrop` to help with this, as well as a new `tcp` bcc library for later processing in Python. I'm sure this code will get reused (and improved) over time.

---

*You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).*