

opensnoop For Linux

25 Jul 2014

Here's another tool I didn't think would be possible: what files are being opened on my Linux system?

```
# ./opensnoop
Tracing open()s. Ctrl-C to end.
COMM      PID      FD FILE
[...]
```

sshd	6503	-1	/usr/share/ssh/blacklist.RSA-2048
sshd	6503	-1	/etc/ssh/blacklist.RSA-2048
sshd	6503	0x4	/root/.ssh/authorized_keys
sshd	6503	-1	/var/run/nologin
sshd	6503	-1	/etc/nologin
sshd	6503	0x4	/etc/passwd
sshd	6503	0x4	/etc/shadow
sshd	6503	0x4	/etc/localtime
sshd	6503	0x4	/etc/security/capability.conf
sshd	6503	0x4	/etc/login.defs

```
[...]
```

This is tracing all processes system-wide, and using a more efficient tracer than [strace](#).

How about just files containing "conf", to see what config files are actively opened?

```
# ./opensnoop conf
Tracing open()s for filenames containing "conf". Ctrl-C to end.
COMM      PID      FD FILE
run       19107    0x3 /etc/nsswitch.conf
stat      19111    0x3 /etc/nsswitch.conf
webapp    19119    0x4 /home/webapp/3.7/rel/conf/logging.conf
webapp    19119    0x7 /etc/nsswitch.conf
webapp    19119    0x7 /etc/host.conf
webapp    19119    0x7 /etc/resolv.conf
webapp    19119    0x8 /home/webapp/3.7/rel/conf/application.properties
catalina.sh 19107    0x3 /etc/nsswitch.conf
run       19122    0x3 /etc/nsswitch.conf
[...]
```

Nice! So that's where my webapp config files are...

How about files ending in "log"?

```
# ./opensnoop 'log$'
Tracing open()s for filenames containing "log$". Ctrl-C to end.
COMM      PID      FD FILE
webapp    21144    0x4 /var/tmp/webapp/3.7/stash/webapp.log
webapp    21159    0x4 /var/tmp/webapp/3.7/stash/webapp.log
webapp    21174    0x4 /var/tmp/webapp/3.7/stash/webapp.log
bash      21301    0x4 /var/log/lastlog
[...]
```

Oh, no wonder I couldn't find them, I wasn't looking in /var/tmp!...

And open()s that returned an error (eg, file not found)?

```
# ./opensnoop -x
Tracing open()s. Ctrl-C to end.
COMM      PID      FD FILE
smtp      2690     -1 /usr/lib/tls/x86_64/libnss_db.so.2
smtp      2690     -1 /usr/lib/tls/libnss_db.so.2
smtp      2690     -1 /usr/lib/x86_64/libnss_db.so.2
smtp      2690     -1 /usr/lib/libnss_db.so.2
java      4680     -1 /home/tomcat/conf/Standalone/localhost
java      4680     -1 /home/tomcat/conf/Standalone/localhost
java      4680     -1 /home/tomcat/conf/Standalone/localhost
[...]
```

Catching a process hunt around its library path is normal, like we see here for smtp. But I wonder if that java app was supposed to be finding that file...

Options

[opensnoop](#) options are summarized by the USAGE message (there's also a [man page](#) and [examples file](#)):

```
# ./opensnoop -h
USAGE: opensnoop [-htx] [-d secs] [-p PID] [-n name] [filename]
      -d seconds      # trace duration, and use buffers
      -n name         # process name to match on I/O issue
      -p PID          # PID to match on I/O issue
      -t              # include time (seconds)
      -x              # only show failed opens
      -h              # this usage message
      filename        # match filename (partials, REs, ok)

eg,
  opensnoop           # watch open()s live (unbuffered)
  opensnoop -d 1      # trace 1 sec (buffered)
  opensnoop -p 181     # trace I/O issued by PID 181 only
  opensnoop conf       # trace filenames containing "conf"
  opensnoop 'log$'     # filenames ending in "log"
```

The -p option employs an in-kernel filter for efficiency.

opensnoop traces events as they happen, which for very frequent open()s can begin to cost measurable overhead. The "-d" mode buffers and prints the buffer at the end, reducing overheads if needed.

What, Why, and How

This is another ftrace-based hack for my [perf-tools](#) collection.

There are some great ways to implement opensnoop on Linux using tracers such as `perf_events` (the "perf" command) with kernel debuginfo, or using SystemTap or ktap. But my aim is to use this now on a fleet of AWS EC2 cloud instances running Linux 3.2, and *without* kernel debuginfo (which is impractical to install everywhere in this dynamic environment, since it can be 100s of Mbytes).

The without-debuginfo requirement makes this especially tricky. One problem was that I didn't think ftrace (or `perf_events`) were able to dereference strings, having only seen examples with hex addresses. It turns out they do have this capability, and in fact, opensnoop can be a few one-liners:

```
# perf probe --add 'do_sys_open filename:string'
[...]
# perf record --no-buffering -e probe:do_sys_open -o - -a | PAGER=cat perf script -i -
multilog 19075 [001] 5586.323974: probe:do_sys_open: (ffffffff811af8b0) filename_string="."
multilog 19075 [001] 5586.324013: probe:do_sys_open: (ffffffff811af8b0) filename_string="/main"
multilog 19075 [001] 5586.324028: probe:do_sys_open: (ffffffff811af8b0) filename_string="lock"
  snmpd 1255 [000] 5586.576142: probe:do_sys_open: (ffffffff811af8b0) filename_string="/proc/r
  snmpd 1255 [000] 5586.576278: probe:do_sys_open: (ffffffff811af8b0) filename_string="/proc/r
[...]
# perf probe --del do_sys_open
```

That's an example of dynamic tracing of the kernel function `do_sys_open()`, and examining an argument (filename) as a string. Except this needs kernel debuginfo to recognize the "filename" symbol.

Note that on older systems (including Linux 3.2), instead of `--no-buffering` it is `-D`.

Without kernel debuginfo, I can fashion a similar one-liner using CPU registers instead of symbols. I was struggling to get strings to work with this, and started thinking it wouldn't be possible, then asked for help on the [perf-users mailing list](#). Fortunately, it is possible!

Instead of using registers on the entry of `do_sys_open()`, I ended up tracing the return value of `getname()` (which `do_sys_open()` calls), and processing that as either a `char *` or `struct filename *` depending on the kernel version. It's a brittle hack, but I thought this was a bit better than guessing the register for `do_sys_open()` on different platforms.

Conclusion

Tracing `open()`s can tell you a lot about running applications: the locations of their config, log, and data files, and file open errors. This is a Linux port of my popular opensnoop tool, which I've used for many years to help with performance analysis and troubleshooting.

This is also another proof of concept for older Linux kernels, like [iosnoop](#), using the existing ftrace and kprobes tracing frameworks. If you happen to have kernel debuginfo, you can also just use the one-liners I included above. Warnings apply: opensnoop and those one-liners use dynamic tracing on Linux, which has had kernel panic bugs in the past, so know what you are doing, test first, and use at your own risk.

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).