

Ubuntu Xenial bcc/BPF

14 Jun 2016

Ubuntu 16.04 LTS (aka Ubuntu Xenial) was released about two months ago, and can run new performance analysis and troubleshooting tools thanks to its Linux 4.4 kernel. These tools use [enhanced BPF](#), and there's a growing number of front ends, including [bcc](#), which I've written about before. Let's take a quick look.

Installation

The following three commands install bcc on Xenial, and are from the [Ubuntu Xenial install docs](#) (UPDATED 3-Oct-2016):

```
# echo "deb [trusted=yes] https://repo.iovisor.org/apt/xenial xenial-nightly main" | sudo tee /etc
deb [trusted=yes] https://repo.iovisor.org/apt/xenial xenial-nightly main
# sudo apt-get update
[...]
# sudo apt-get install bcc-tools
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  binutils libbcc libpython-stdlib libpython2.7-minimal libpython2.7-stdlib python python-bcc pyth
Suggested packages:
  binutils-doc python-doc python-tk python2.7-doc binfmt-support
The following NEW packages will be installed:
  bcc-tools binutils libbcc libpython-stdlib libpython2.7-minimal libpython2.7-stdlib python pyth
0 upgraded, 11 newly installed, 0 to remove and 80 not upgraded.
Need to get 17.2 MB of archives.
After this operation, 69.3 MB of additional disk space will be used.
Do you want to continue? [Y/n] y
[...]
```

At some point bcc-tools should also be available as an Ubuntu snap, making the install process even easier. The bcc install docs also contain instructions for an [Ubuntu source](#) build of bcc, which you can do if you want to develop your own bcc tools.

Screenshots

Tools appear under /usr/share/bcc/tools. Eg, tracing disk I/O:

```
# /usr/share/bcc/tools/biosnoop
TIME(s)      COMM      PID    DISK    T    SECTOR    BYTES    LAT(ms)
0.000000000  bash      2952   sda     R    9457088    4096     0.99
0.000979000  bash      2952   sda     R    9457072    4096     0.48
0.001760000  bash      2952   sda     R    9483416    4096     0.43
0.006208000  bash      2967   sda     R    9667584    16384    0.39
0.007636000  cksum     2967   sda     R    9667616    16384    0.66
0.009509000  cksum     2967   sda     R    9664280    53248    0.92
0.010854000  cksum     2967   sda     R    6119520    4096     0.41
0.011985000  cksum     2967   sda     R    9661832    4096     0.35
0.013343000  cksum     2967   sda     R    5887144    24576    0.75
0.014787000  cksum     2967   sda     R    6047760    16384    0.54
[...]
```

This is more efficient than my earlier versions of this tool, as it times events in kernel context, and only sends the summary lines to user space for printing.

Perhaps that's overkill, and you just want to see the latency distribution. Here's biolatility:

```
# /usr/share/bcc/tools/biolatency
Tracing block device I/O... Hit Ctrl-C to end.
^C
      usecs      : count      distribution
      0 -> 1      : 0
      2 -> 3      : 0
      4 -> 7      : 0
      8 -> 15     : 0
     16 -> 31     : 2
     32 -> 63     : 1
     64 -> 127    : 4
    128 -> 255    : 8
    256 -> 511    : 86
    512 -> 1023   : 94
   1024 -> 2047   : 43
   2048 -> 4095   : 1
                                     *
                                     ***
                                     *****
                                     *****
                                     *****
```

This tool uses BPF to create a custom histogram that is counted in kernel context, for efficiency. Only an array of numbers (the "count" column) is transferred to user space.

These tools usually have Unix-like usage, eg:

```
# cd /usr/share/bcc/tools
# ./biolatility -h
usage: biolatility [-h] [-T] [-Q] [-m] [-D] [interval] [count]

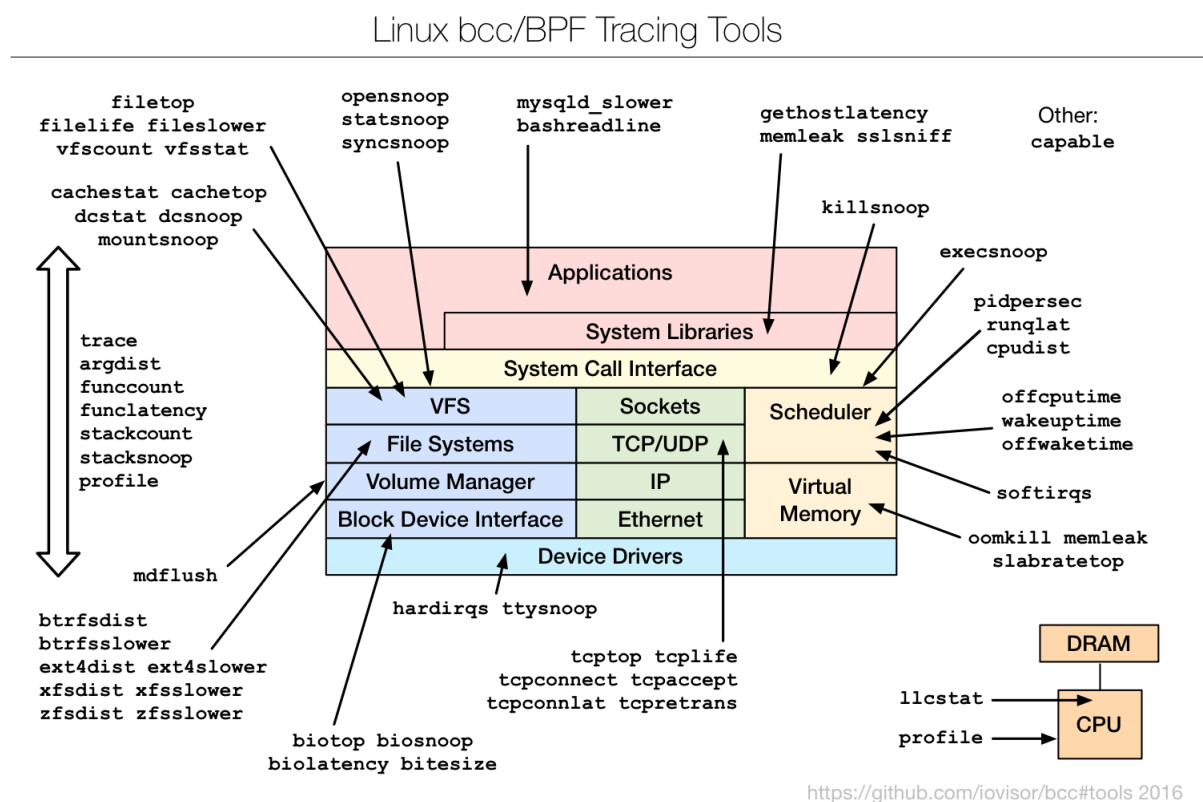
Summarize block device I/O latency as a histogram

positional arguments:
  interval      output interval, in seconds
  count         number of outputs

optional arguments:
  -h, --help            show this help message and exit
  -T, --timestamp       include timestamp on output
  -Q, --queued          include OS queued time in I/O time
  -m, --milliseconds   millisecond histogram
  -D, --disks           print a histogram per disk device

examples:
./biolatility          # summarize block I/O latency as a histogram
./biolatility 1 10     # print 1 second summaries, 10 times
./biolatility -mT 1    # 1s summaries, milliseconds, and timestamps
./biolatility -Q       # include OS queued time in I/O time
./biolatility -D       # show each disk device separately
```

There are lots more tools so far, and more will be added. Here's a diagram:



Finally, one more of my favorites, `ext4slower`, which can show I/O slower than a custom threshold, timed and filtered in kernel context:

```
# ./ext4slower 1
Tracing ext4 operations slower than 1 ms
TIME      COMM      PID    T BYTES  OFF_KB  LAT(ms)  FILENAME
03:36:05  cksum      2738   R 29162   0        1.32     h2ph
03:36:05  cksum      2738   R 27280   0        1.39     hostid
03:36:05  cksum      2738   R 65536   0        1.15     info
03:36:05  cksum      2738   R 65536   0        1.01     lexgrog
03:36:05  cksum      2738   R 1577    0        2.10     linux-boot-prober
03:36:05  cksum      2738   R 65536   0        1.14     localectl
03:36:05  cksum      2738   R 65536   128      1.48     localectl
03:36:05  cksum      2738   R 65536   0        1.04     ltrace
03:36:05  cksum      2738   R 65536   128      1.84     ltrace
03:36:05  cksum      2738   R 65536   640      1.25     lxc
03:36:06  cksum      2738   R 65536   11008    5.37     lxd
```

You can browse the `*example.txt` files in the [bcc tools directory](#) for more screenshots and examples. Some are single-purpose tools, like [biolatility](#) and [ext4slower](#), and some are multi-tools, like [trace](#) and [argdist](#).

Documentation

All tools have man pages, and example files (see previous links).

```
# export MANPATH=$MANPATH:/usr/share/bcc/man
# man biosnoop
biosnoop(8)                                System Manager's Manual                biosnoop(8)

NAME
    biosnoop - Trace block device I/O and print details incl. issuing PID.

SYNOPSIS
    biosnoop

DESCRIPTION
    This tools traces block device I/O (disk I/O), and prints a one-line
    summary for each I/O showing various details. These include the latency
    from the time of issue to the device to its completion, and the PID and
    process name from when the I/O was first created (which usually identi-
    fies the responsible process).

[...]
```

Caveats

Usual warnings about tracers apply: even though BPF is JIT-optimized kernel context summaries, you could trace such a high frequency of events (millions/sec) that the performance overhead slows other apps. Test in a lab environment before use.

Ubuntu Xenial is also a 4.4 kernel, and the bcc tools are developed for the latest kernel, so some tools may not work. We've been putting older versions under /usr/share/bcc/tools/old, which might be worth checking. For example, tools that print stack traces (eg, stackcount) won't work on Xenial (which is missing the newer BPF stack trace support), but the tools/old versions may work instead (although they were workarounds until proper support existed).

bcc/BPF has more capabilities that the tools don't use yet, but will over time. See my [bcc/BPF road ahead](#) post for what's next.

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).