

USENIX/LISA 2014 New Tools and Old Secrets (perf-tools)

17 Mar 2015

At the last USENIX/LISA conference, I gave a talk on new Linux performance tools: my open source [perf-tools](#) collection. These use existing kernel frameworks, ftrace and perf_events, which are built in to most Linux kernel distributions by default, including the Linux cloud instances I analyze at Netflix.

The slides are on [slideshare](#):



The talk was also videoed, which is on [youtube](#):

Linux Performance Analysis: New Tools and Old Secrets



I have a long history of creating DTrace scripts, published in the DTraceToolkit and the DTrace book, some of which are shipped by default in some OSes. After switching to Linux, I've been looking for ways to port them over, and have found that ftrace and perf_events – which are in the Linux kernel source – can provide many of the capabilities I need. My perf-tools collection provides scripts to facilitate their use.

Ftrace, in particular, seems to be a well-kept secret of the Linux kernel. It was created by Steven Rostedt, and has had virtually no marketing. The tools I'm creating can help raise awareness, by showing examples of what ftrace can do.

For example, my ftrace-based iosnoop:

```
# ./iosnoop
Tracing block I/O... Ctrl-C to end.
COMM      PID    TYPE DEV    BLOCK      BYTES      LATms
supervise 1809   W   202,1  17039968   4096       1.32
supervise 1809   W   202,1  17039976   4096       1.30
tar       14794  RM   202,1  8457608    4096       7.53
tar       14794  RM   202,1  8470336    4096      14.90
tar       14794  RM   202,1  8470368    4096       0.27
tar       14794  RM   202,1  8470784    4096       7.74
tar       14794  RM   202,1  8470360    4096       0.25
tar       14794  RM   202,1  8469968    4096       0.24
tar       14794  RM   202,1  8470240    4096       0.24
[...]
```

The output shows the tar command making a series of metadata reads (type == RM), with one taking 14.9 milliseconds. Like tcpdump, but for disks, this can be useful to track down odd performance behaviors that may involve sequences of I/O and their queueing effects.

The perf-tools collection contains several single-purpose tools, like iosnoop, which aim to do one thing and do it well (Unix philosophy).

There are also multi-purpose tools, which require more expertise to use, but are also more powerful. For example, I'll use funccount to count kernel calls beginning with "ip":

```
# ./funccount 'ip*'
Tracing "ip*"... Ctrl-C to end.
^C
FUNC                                COUNT
[...]
ip_mc_sf_allow                      70
ipv6_chk_mcast_addr                 72
ip_finish_output                    108
ip_local_out                        108
ip_output                           108
ip_queue_xmit                       108
ipv4_mtu                            216
ip_local_deliver                    229
ip_local_deliver_finish             229
ip_rcv                              229
ip_rcv_finish                       229
ipv4_dst_check                      513

Ending tracing...
```

This output shows 108 calls to `ip_output()`, which sounds like a function for sending IP packets. Let's pull out some stack traces to see how we got here, using `kprobe`:

```
# ./kprobe -s 'p:ip_output' | head -50 > out.ip_output
# more out.ip_output
Tracing kprobe ip_output. Ctrl-C to end.
      sshd-19389 [003] d... 3042954.165578: ip_output: (ip_output+0x0/0x90)
      sshd-19389 [003] d... 3042954.165584:
=> ip_queue_xmit
=> tcp_transmit_skb
=> tcp_write_xmit
=> __tcp_push_pending_frames
=> tcp_sendmsg
=> inet_sendmsg
=> sock_aio_write
=> do_sync_write
=> vfs_write
=> Sys_write
=> system_call_fastpath
      sshd-19250 [001] d... 3042954.258718: ip_output: (ip_output+0x0/0x90)
      sshd-19250 [001] d... 3042954.258723:
=> ip_queue_xmit
[...]
```

Nice! We can see the path from the `write()` syscall to `ip_output()`, through VFS and TCP. I redirected the output of `kprobe` to a file to avoid a feedback loop, as I'm logged in via SSH. The `head` command ensures that `kprobe` exits after capturing several stack traces, as we don't want to leave this running (overhead).

I can go further with `kprobe`, and inspect functions arguments as well. There are [examples of kprobe](#) and all other tools in the `perf-tools` collection. If need be, I can re-instrument the same `kprobe` one-liner using `perf_events`, which has lower overhead.

While these tools have a polished interface (USAGE message, man page, examples file), their internals often make use of temporary hacks, awaiting more Linux kernel features. For example, `iosnoop` passes both I/O request and response timestamps to user-level, which then calculates the delta, when it would be more efficient to do this calculation in-kernel, and only pass the result.

In particular, I'm looking forward to eBPF (or a similar facility) being available in the kernel, so that I can do more kernel-level programming including the I/O latency calculation during tracepoints, and reduce the overhead of tracing. I'll update my tools to use eBPF when it is available, and I'll also be able to create more.

LISA is a great conference, although I arrived late as there was a clash with AWS re:Invent, which I also spoke at. I can at least see the talks I missed, since they were videoed: they are linked on the [conference program](#).

I did arrive in time to see some talks, including Ben Rockwood's excellent [I Am SysAdmin \(And So Can You!\)](#), which I referenced. Despite the many hats I've worn, I still feel like a sysadmin, as I still use those skills day to day. Ben and I also took a moment to pose with Deirdré, who created the ponycorn mascots in my slidedeck, for a [photo](#). I hope to be back at LISA in 2015.

You can comment here, but I can't guarantee your comment will remain here forever: I might switch comment systems at some point (eg, if Disqus add advertisements).

Copyright 2017 Brendan Gregg.

[About this blog](#)

[Perf Methods](#)

[USE Method](#)

[TSA Method](#)

[Off-CPU Analysis](#)

[Active Bench.](#)

[Flame Graphs](#)

[Heat Maps](#)

[Frequency Trails](#)

[Colony Graphs](#)

[perf Examples](#)

[eBPF Tools](#)

[DTrace Tools](#)

[DTraceToolkit](#)

[DtkshDemos](#)

[Guessing Game](#)

[Specials](#)

[Books](#)

[Other Sites](#)